

# Parallel Gaussian Elimination



# Gaussian Elimination

- We work with the rowwise decomposition:
  - Each processor receives  $n/p$  rows of the extended matrix  $(A,b)$
- We parallelize the individual pivoting steps
  - In step  $i$  we find the maximum  $A(j,i)$  for  $i \leq j \leq n$
  - Swap row  $j$  with row  $i$
  - Broadcast row  $i$  to processors holding rows  $i+1 \dots n$
  - Each processor computes the new  $\text{row}_j = \text{row}_j - A(j,i)/A(i,i) \cdot \text{row}_i$

# Gaussian Elimination

- Time: Finding the maximum  $A(j,i)$  and  $j$  can be done by tournaments in time  $O(n/p + \log p)$
- Broadcasting the pivot row takes communication time  $O(n \cdot \log p)$
- Computation time for the remaining operations is  $O(n^2/p)$  per phase
- In  $n$  phases the computation time is thus  $O(n^3/p)$  and the communication time is  $O(n^2 \log p)$
- Hence the algorithm has constant efficiency once  $n = \Omega(p \log p)$

# Gaussian Elimination

- The disadvantage of this algorithm is that during the broadcasting step no computations take place, so effectively processors are idle during these communications, for time  $O(n^2p)$

# A different algorithm

- Again we work with the rowwise decomposition
- In the previous algorithm we were using tournaments to find the maximum  $A(j,i)$
- Consequently we couldn't predict the pivot row  $j$  ahead of time
- And must then broadcast that row

# A different algorithm

- Again we parallelize the pivoting step
- This time we always use row  $i$  in phase  $i$
- Find the maximum  $A(i,j)$
- We then want to eliminate all  $A(k,j)$  for  $k > i$
- When we do this for  $i=1, \dots, n$  the resulting matrix is a columnwise permuted triangular matrix
- So we can afterwards solve via backsubstitution

# A different algorithm

- The big advantage is that now we can organize computation and communication more efficiently
  - Processors are arranged as a linear array
  - We know that row  $i$  is the pivoting row
  - The processor holding row  $i$  sends row  $i$  to its nearest neighbor in the array
  - Every processor immediately passes the row on after reception
  - Then starts to replace  $\text{row}_k = \text{row}_k - A(k,j)/A(i,j) \cdot \text{row}_i$  for all rows  $k$  in its possession

# A different algorithm

- Time: The computation time is again  $O(n^3/p)$
- The communication time:
  - The last processor can start working after  $(p-1)n$  communication time
  - However, this delay happens only once, when the computation time  $n^2/p$  per phase exceeds  $pn$ , i.e., if  $n = \Omega(p^2)$ , because
  - In that case the overall extra communication time is only  $O(pn)$
  - This effect is called pipelining